

IRGAN: A Minimax Game for Unifying Generative and Discriminative Information Retrieval Models

Jun Wang
University College London
j.wang@cs.ucl.ac.uk

Lantao Yu, Weinan Zhang^{*}
Shanghai Jiao Tong University
wnzhang@sjtu.edu.cn

Yu Gong, Yinghui Xu
Alibaba Group
renji.xyh@taobao.com

Benyou Wang, Peng Zhang
Tianjin University
pzhang@tju.edu.cn


Dell Zhang
Birkbeck, University of London
dell.z@ieee.org

2017 SIGIR

王润生
May 5, 2019

The Idea of GANs

- There is a generator **G** that captures the data distribution and tries to generate high quality data to fool the discriminator **D**, to maximize the probability of **D** making a mistake
- There is a discriminator **D** that estimates the probability that a sample came from the real data set rather than **G** and tries to distinguish them



The Formulation of GANs

$$\min_G \max_D V(D, G)$$

$$V(D, G) = \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{z \sim p_z(z)} [\log (1 - D(G(z)))]$$

$$= \mathbb{E}_{x \sim p_{data}(x)} [\log D(x)] + \mathbb{E}_{x \sim p_g(x)} [\log (1 - D(x))]$$

For G fixed, the optimal D is

$$D^* = \max_D V = \max_D \int_x p_{data}(x) \log(D(x)) + p_g(x) \log(1 - D(x)) dx$$

$$= \max_D a * \log(D) + b * \log(1 - D)$$

$$= \frac{a}{a+b}, \text{ where } a = p_{data}(x), b = p_g(x)$$

The Formulation of GANs

For D fixed, the optimal G is

$$\begin{aligned} G^* &= \min_G \int_x p_{data}(x) \log(D^*(x)) + p_g(x) \log(1 - D^*(x)) dx \\ &= \min_G \int_x p_{data}(x) \log \frac{p_{data}(x)}{p_{data}(x) + p_g(x)} dx \\ &\quad + \int_x p_g(x) \log \left(\frac{p_g(x)}{p_{data}(x) + p_g(x)} \right) dx \\ &= \min_G \log \frac{1}{2} + D_{KL}(P_{data} \| \mathbf{M}) + \log \frac{1}{2} + D_{KL}(P_g \| \mathbf{M}) \\ &= \min_G -\log 4 + 2 D_{JS}(P_{data} \| P_g), \text{ where } \mathbf{M} = \frac{p_{data} + p_g}{2} \end{aligned}$$

The Training of GANs

- Training Discriminator

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log (1 - D(G(z^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(z^{(i)}))).$$

end for

The Training of GANs

- Training Generator

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Sample minibatch of m examples $\{x^{(1)}, \dots, x^{(m)}\}$ from data generating distribution $p_{\text{data}}(x)$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(x^{(i)}) + \log \left(1 - D(G(z^{(i)})) \right) \right].$$

end for


- Sample minibatch of m noise samples $\{z^{(1)}, \dots, z^{(m)}\}$ from noise prior $p_g(z)$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log \left(1 - D(G(z^{(i)})) \right).$$


end for

Two Schools of Thinking in Information Retrieval Modeling

- Generative Retrieval
 - Focus on predicting relevant documents given a query



- Discriminative Retrieval
 - Focus on predicting relevancy given a query-document pair



Unify Two Schools of Thinking Based on the Idea of GANs

- Generative Retrieval
 - The generative retrieval model can learn to fit the underlying relevance distribution over documents with the help of discriminator
- Discriminative Retrieval
 - The discriminative retrieval model can not only exploit the true data but also exploit the fake data selected from the generative model to get a better estimation

The Formulation of IRGAN

- Problem Definition

- Given a set of queries $\{q_1, q_2, \dots, q_N\}$ and a set of documents $\{d_1, d_2, \dots, d_M\}$. A query here can be any specific form of the user's information need such as search keywords, a user profile, or a question
- The underlying true relevance distribution $p_{true}(d|q, r)$ describes the user's preference distribution over the candidate documents with respect to her submitted query

The Formulation of IRGAN

- Problem Definition
 - Training set
 - Given a set of samples from $p_{true}(d|q, r)$ as the training data
 - Generative Retrieval Model
 - $p_{\theta}(d|q, r)$ which generates/selects relevant documents from the candidate pool for the query q , trying to approximate the true relevance distribution
 - Discriminative Retrieval Model
 - $f_{\phi}(q, d)$ which tries to distinguish between true and fake

The Formulation of IRGAN

- Overall Objective

$$J^{G^*, D^*} = \min_{\theta} \max_{\phi} \sum_{n=1}^N \left(\mathbb{E}_{d \sim p_{\text{true}}(d|q_n, r)} [\log D(d|q_n)] + \mathbb{E}_{d \sim p_{\theta}(d|q_n, r)} [\log(1 - D(d|q_n))] \right)$$

where

$$D(d|q) = \sigma(f_{\phi}(d, q)) = \frac{\exp(f_{\phi}(d, q))}{1 + \exp(f_{\phi}(d, q))}$$

The Formulation of IRGAN

- Optimizing Discriminative Retrieval

$$\phi^* = \arg \max_{\phi} \sum_{n=1}^N \left(\mathbb{E}_{d \sim p_{\text{true}}(d|q_n, r)} \left[\log(\sigma(f_{\phi}(d, q_n))) \right] + \mathbb{E}_{d \sim p_{\theta^*}(d|q_n, r)} \left[\log(1 - \sigma(f_{\phi}(d, q_n))) \right] \right)$$

The Formulation of IRGAN

- Optimizing Generative Retrieval

$$\begin{aligned}\theta^* &= \arg \min_{\theta} \sum_{n=1}^N \left(\mathbb{E}_{d \sim p_{\text{true}}(d|q_n, r)} \left[\log \sigma(f_{\phi}(d, q_n)) \right] + \right. \\ &\quad \left. \mathbb{E}_{d \sim p_{\theta}(d|q_n, r)} \left[\log(1 - \sigma(f_{\phi}(d, q_n))) \right] \right) \\ &= \arg \max_{\theta} \sum_{n=1}^N \underbrace{\mathbb{E}_{d \sim p_{\theta}(d|q_n, r)} \left[\log(1 + \exp(f_{\phi}(d, q_n))) \right]}_{\text{denoted as } J^G(q_n)}\end{aligned}$$

$$\sigma(f_{\phi}(d, q)) = \frac{\exp(f_{\phi}(d, q))}{1 + \exp(f_{\phi}(d, q))}$$

The Formulation of IRGAN

- Optimizing Generative Retrieval
 - As the sampling of d is discrete, it cannot be directly optimized by SGD from the original GANs
 - 离散数据无法直接反向传播，一般采用 RL 中的 policy gradient 或者 gumbel softmax 来近似计算
 - 离散数据和连续数据 embedding 空间之间存在差异，embedding 上的一点微小改变，很难直接反映到离散数据上
 - The view of Ian Goodfellow:

The Formulation of IRGAN

- The View of Ian Goodfellow:
 - You can make slight changes to the synthetic(合成) data only if it is based on continuous numbers. If it is based on discrete numbers, there is no way to make a slight change.
 - For example, if you output an image with a pixel value of 1.0, you can change that pixel value to 1.0001 on the next step.
 - If you output the word “penguin”(企鹅), you can’t change that to “penguin + .001” on the next step, because there is no such word as “penguin + .001”. You have to go all the way from “penguin” to “ostrich”(鸵鸟).
- Policy gradient based reinforcement learning is used here

The Formulation of IRGAN

- Optimizing Generative Retrieval

$$\nabla_{\theta} J^G(q_n)$$

$$= \nabla_{\theta} \mathbb{E}_{d \sim p_{\theta}(d|q_n, r)} [\log(1 + \exp(f_{\phi}(d, q_n)))]$$

$$\nabla_{\theta} \log(p_{\theta}) = \frac{1}{p_{\theta}} \nabla_{\theta} p_{\theta}$$

$$= \sum_{i=1}^M \nabla_{\theta} p_{\theta}(d_i|q_n, r) \log(1 + \exp(f_{\phi}(d_i, q_n)))$$

$$= \sum_{i=1}^M p_{\theta}(d_i|q_n, r) \nabla_{\theta} \log p_{\theta}(d_i|q_n, r) \log(1 + \exp(f_{\phi}(d_i, q_n)))$$

$$= \mathbb{E}_{d \sim p_{\theta}(d|q_n, r)} [\nabla_{\theta} \log p_{\theta}(d|q_n, r) \log(1 + \exp(f_{\phi}(d, q_n)))]$$

$$\simeq \frac{1}{K} \sum_{k=1}^K \nabla_{\theta} \log p_{\theta}(d_k|q_n, r) \log(1 + \exp(f_{\phi}(d_k, q_n)))$$

Reward Term


Demonstrate of Policy Gradient

- $G(d) = P_{\theta}(d|q)$, where $d \sim \{d_1, d_2, d_3, d_4, d_5\}$

- $$P(d_i) = \frac{\exp(f(d_i))/\tau}{\sum_j \exp(f(d_j))/\tau}$$


where f is the function with respect to specific task

- The Relevance Distribution of q




Demonstrate of Policy Gradient

- Step 1: initialize θ



- Step 2: Generate d5
Observe *negative* reward from D


- Step 3: update θ by policy gradient



- Step 4: Generate d1
Observe *positive* reward from D


Demonstrate of Policy Gradient

- Step 7:update θ by policy gradient



- Step 8.....

- Step 5:update θ by policy gradient



- Step 6:Generate d1
Observe *positive* reward from D

The Algorithm of IRGAN

Algorithm 1 Minimax Game for IR (a.k.a IRGAN)

Input: generator $p_{\theta}(d|q, r)$; discriminator $f_{\phi}(\mathbf{x}_i^q)$;
training dataset $\mathcal{S} = \{\mathbf{x}\}$

- 1: Initialise $p_{\theta}(d|q, r), f_{\phi}(q, d)$ with random weights θ, ϕ .
 - 2: Pre-train $p_{\theta}(d|q, r), f_{\phi}(q, d)$ using \mathcal{S}
 - 3: **repeat**
 - 4: **for** g-steps **do**
 - 5: $p_{\theta}(d|q, r)$ generates K documents for each query q
 - 6: Update generator parameters via policy gradient Eq. (5)
 - 7: **end for**
 - 8: **for** d-steps **do**
 - 9: Use current $p_{\theta}(d|q, r)$ to generate negative examples and combine with given positive examples \mathcal{S}
 - 10: Train discriminator $f_{\phi}(q, d)$ by Eq. (3)
 - 11: **end for**
 - 12: **until** IRGAN converges
-

Application on Item Recommendation

$$p_{\theta}(d|q, r) = \frac{\exp(g_{\theta}(q, d)/\tau)}{\sum_{j \in \mathcal{I}} \exp(g_{\theta}(q, j)/\tau)}$$

$$D(d|q) = \sigma(f_{\phi}(d, q)) = \frac{\exp(f_{\phi}(d, q))}{1 + \exp(f_{\phi}(d, q))}$$

$$g_{\theta}(q, d) = s_{\theta}(q, d) \quad \text{and} \quad f_{\phi}(q, d) = s_{\phi}(q, d)$$

$$s(u, i) = b_i + \mathbf{v}_u^{\top} \mathbf{v}_i$$

Experiment to Item Recommendation

Table 3: Item recommendation results (Movielens).

	P@3	P@5	P@10	MAP
MLE	0.3369	0.3013	0.2559	0.2005
BPR [35]	0.3289	0.3044	0.2656	0.2009
LambdaFM [45]	0.3845	0.3474	0.2967	0.2222
IRGAN-pointwise	0.4072	0.3750	0.3140	0.2418
Impv-pointwise	5.90%*	7.94%*	5.83%*	8.82%*
	NDCG@3	NDCG@5	NDCG@10	MRR
MLE	0.3461	0.3236	0.3017	0.5264
BPR [35]	0.3410	0.3245	0.3076	0.5290
LambdaFM [45]	0.3986	0.3749	0.3518	0.5797
IRGAN-pointwise	0.4222	0.4009	0.3723	0.6082
Impv-pointwise	5.92%*	6.94%*	5.83%*	4.92%*

Table 4: Item recommendation results (Netflix).

	P@3	P@5	P@10	MAP
MLE	0.2941	0.2945	0.2777	0.0957
BPR [35]	0.3040	0.2933	0.2774	0.0935
LambdaFM [45]	0.3901	0.3790	0.3489	0.1672
IRGAN-pointwise	0.4456	0.4335	0.3923	0.1720
Impv-pointwise	14.23%*	14.38%*	12.44%*	2.87%*
	NDCG@3	NDCG@5	NDCG@10	MRR
MLE	0.3032	0.3011	0.2878	0.5085
BPR [35]	0.3077	0.2993	0.2866	0.5040
LambdaFM [45]	0.3942	0.3854	0.3624	0.5857
IRGAN-pointwise	0.4498	0.4404	0.4097	0.6371
Impv-pointwise	14.10%*	14.27%*	13.05%*	8.78%*

The End

Thanks for Listening